NEURAI PLATFORM

Executive demo — speaker script

A dedicated, timed narrative for each of the 10 presentation slides

Repository-grounded final edition · 26 August 2026

Architecture, product, trust, speech intelligence, market position, and expansion thesis

PRESENTATION MAP

A 13-minute story with room for a live product moment

Problem → wedge → product → AI-native model → architecture → trust → Persian speech → competition → expansion

Slide	Claim	Time
1	The AI-native operating surface for organizational memory	0:00–1:10
2	Truth is created in conversation and disappears into fragments	1:10–2:20
3	One conversation can become truth, memory, reasoning, and action	2:20–3:35
4	The platform hub and Echo feel like one system—not a bundle of AI features	3:35–5:00
5	The assistant moves from chat to a governed operating layer	5:00–6:15
6	Experience, control, and data scale independently without losing identity	6:15–7:50
7	JWT, identity, RLS, grants, policy, and consent form one wall	7:50–9:10
8	Accuracy is a pipeline; trust survives when fidelity changes	9:10–10:35
9	Gong, Otter, and Fireflies validate demand—NeurAI changes the axis	10:35–12:05
10	One governed memory substrate can power an entire AI-native platform	12:05–13:20

> **Delivery style** Excited, precise, and credible. Treat every diagram as a story, not a checklist. Keep the strongest words—memory, authority, Persian-first, presence—consistent across the talk.

SLIDE 1 · 0:00–1:10

NeurAI Platform

The AI-native operating surface for organizational memory

> **Purpose** Open with the platform ambition, then make trust and Persian-first design part of the promise—not footnotes.

Suggested speech

What if AI did not begin when we opened a chat? What if it was already present where the work happened—listening with consent, remembering with provenance, and acting only with our authority? That is NeurAI Platform.

We are building an AI-native operating surface for organizations. It captures the truth of work, makes that truth usable across time, and gives people an assistant that can reason and act without becoming a security exception.

Our first app is Echo. Echo turns calls and meetings into transcripts, speakers, versioned summaries, searchable memory, and governed actions. Persian is not a translated skin here. The interface, speech pipeline, search behavior, names, dates, digits, and direction all begin Persian-first.

The promise is simple: capture what happened, explain what it means, act only with authority, and learn across time.

Presenter cue

Pause after ‘That is NeurAI Platform.’ Let the cover image establish scale before naming Echo.

Transition

To understand why this matters, start with the loss every organization already accepts as normal.

SLIDE 2 · 1:10–2:20

The problem is not more meetings—it is memory loss

Truth is created in conversation and disappears into fragments

> **Purpose** Turn meeting pain into a strategic data problem.

Suggested speech

Organizations do not suffer because they lack information. They suffer because the most important information is born in conversation and then decays.

A commitment lives in someone's notes. A customer objection stays in a recording. A correction is repeated three meetings later. Context is retyped into a generic assistant, usually without source, scope, or history.

Traditional meeting tools stop at notes. Traditional enterprise systems start after the human has entered structured data. The gap between those two moments is where decisions, trust, and time are lost.

NeurAI Platform treats this as a memory architecture problem. We need a source of truth, durable provenance, identity-aware access, and a path from evidence to action.

Presenter cue

Point to the four losses: source, ownership, continuity, follow-through.

Transition

That leads to the most important product decision we made: where to begin.

SLIDE 3 · 2:20–3:35

Echo is the wedge—and the memory substrate

One conversation can become truth, memory, reasoning, and action

Purpose Explain the why-now and why-Echo strategy.

Suggested speech

We start with Echo because conversations are high-frequency, information-dense, and expensive to forget. The value is immediate: record, transcribe, identify speakers, correct, summarize, search, and share.

But the strategic value compounds. The transcript becomes durable memory. The next call can be understood against previous calls. The assistant can cite the moment a decision was made. A workflow can react to a processed call without scraping a dashboard.

Echo also forces us to solve the hard platform problems early: recording consent, sensitive content, multi-tenant isolation, background processing, model provenance, deletion, human correction, and bounded action.

A horizontal assistant without a truth source is generic. A meeting tool without a platform is bounded. Echo gives us both the wedge and the substrate.

Presenter cue

Walk left to right through Capture → Truth → Memory → Reason → Act.

Transition

Now let’s make that concrete in the experience a user sees.

SLIDE 4 · 3:35–5:00

One calm product surface

The platform hub and Echo feel like one system—not a bundle of AI features

Purpose Demo the current experience through two screenshots.

Suggested speech

This is the Persian-first platform hub. The user can ask across permitted memory, attach calls as sources, select an agent or model, and choose whether web search is part of the turn.

Then Echo moves to the moment of truth creation. The recorder keeps waveform, live transcript, agenda, and notes in one place. Audio is buffered locally in the browser, uploaded in bounded parts, and processed without exposing speech-provider credentials.

After processing, the record page keeps audio, transcript, speaker roster, notes, summary versions, and the contextual assistant together. The user never has to choose between evidence and intelligence.

The visual language is intentionally calm: dark neutral surfaces, precise violet emphasis, strong hierarchy, and RTL behavior that feels native rather than mirrored after the fact.

Presenter cue

Show the hub first; then the recorder. If live product is available, use the screenshots as the fallback path.

Transition

The visible product is calm because the intelligence underneath has a very explicit shape.

SLIDE 5 · 5:00–6:15

AI-native means presence + hands + signals

The assistant moves from chat to a governed operating layer

> **Purpose** Explain the platform's AI-native primitives in memorable language.

Suggested speech

Our definition of AI-native has three active layers and one control.

Presence means the assistant is available in the hub and beside the record where context already exists. Hands means it can use tools: search transcripts, read related calls, navigate the product, manage records, or propose a correction. Signals mean important events can create a brief or trigger a governed rule without waiting for a fresh prompt.

The control is autonomy: Watch, Assist, Act. Watch observes and answers. Assist proposes and asks before writes. Act is reserved for pre-approved write classes and remains fully audited.

This is a crucial distinction: autonomy changes consequences, not intelligence. We do not make the model 'more trusted.' We make the exposed effects more explicit.

Presenter cue

Emphasize the words Presence, Hands, Signals, Autonomy. They should become the audience's vocabulary.

Transition

Those primitives work because the architecture gives every kind of authority a place.

SLIDE 6 · 6:15–7:50

Four parts. Three planes. Explicit boundaries.

Experience, control, and data scale independently without losing identity

> **Purpose** Give technical confidence without drowning the room in boxes.

Suggested speech

The platform has four deployable parts. The web application is the experience plane. Its Backend for Frontend keeps secure cookies and internal topology away from the browser. Core is the control plane: the Fastify API, worker, agent runtime, policies, gateway, and live relay. ML is a stateless speech facade. Supabase is the durable data plane: Postgres, Auth, Storage, row-level security, grants, and queues.

The boundaries are about authority. The browser holds a session, not provider secrets. ML gets a short-lived signed URL to one audio object, not a platform credential. Background jobs re-resolve the call owner. Agent tools use a narrower database role.

We chose a small number of real boundaries instead of a large number of nominal services. That keeps the system operable while making the trust seams reviewable.

Presenter cue

Trace one request: Browser → BFF → Core → identity-bound database. Then trace one audio job through Worker → signed URL → ML.

Transition

Now we can show the differentiator that is hardest to copy with a feature checklist: trust as architecture.

SLIDE 7 · 7:50–9:10

The agent borrows authority—never invents it

JWT, identity, RLS, grants, policy, and consent form one wall

Purpose Make the security model understandable and distinctive.

Suggested speech

A JWT is the signed token that tells us who the caller claims to be. We verify it, resolve an active platform user, and attach that identity inside the same database transaction as every query.

Postgres row-level security limits which rows that person can see. Role grants set the ceiling on what the application and agent can do. The agent role has no DELETE. A central tool policy checks declared tools, admin gates, and attempt budgets. Consequential content changes become proposals that a person confirms.

No single layer carries the promise. If UI code is wrong, RLS still applies. If a prompt is manipulated, grants still apply. If a model keeps trying, the tool budget ends the loop.

The profound choice is that the agent has no omniscient service identity. It borrows the caller's authority and never more.

Presenter cue

Build the wall bottom-up. End on 'never more' and pause.

Transition

That same honesty appears in the most technically demanding part of Echo: Persian speech.

SLIDE 8 · 9:10–10:35

Persian-first speech—with honest degradation

Accuracy is a pipeline; trust survives when fidelity changes

Purpose Show product and engineering depth around speech.

Suggested speech

Persian-first is not a language toggle. It shapes the UI, text normalization, digits, dates, mixed-script typography, search, speech hints, and the organization's glossary of names and terms.

The audio pipeline measures channels, avoids the dual-mono trap, runs local voice activity detection, uses Soniox as the primary STT lane with an OpenRouter fallback, and adds local diarization where a mono stream needs speaker structure. Voice enrollment can conservatively connect recurring speakers to directory people.

Most importantly, fidelity degrades honestly. Word timestamps give click-a-word. Segment timestamps give click-a-line. If a provider gives prose with no timing, Echo anchors it to a real speech span instead of returning zeros or nothing.

A correct transcript that cannot be traced back to audio is not complete. Our timing ladder keeps evidence usable.

Presenter cue

Use the timing ladder as the memorable technical example: word → line → anchored span.

Transition

With that product and architecture in mind, how do we position against a proven category?

SLIDE 9 · 10:35–12:05

The category is proven; our opening is distinct

Gong, Otter, and Fireflies validate demand—NeurAI changes the axis

> **Purpose** Frame competition with respect and a specific wedge.

Suggested speech

Gong proves that conversation intelligence can become an operating system for revenue. Otter proves that general meeting knowledge can become conversational. Fireflies proves the pull for integrations, skills, and meeting automation.

We should respect those strengths. Our opportunity is not to copy every checklist. NeurAI Platform is differentiated by three connected choices: Persian-first depth, authority-aware agents, and a platform designed to grow beyond meeting notes.

Otter’s current official language list does not include Persian. Fireflies and Gong are not publicly positioned as Persian-first operating systems. More importantly, our agent, worker, and database all carry the same caller-bound authority model.

Our strategy is to win an underserved trust-and-language surface, then compound Echo’s memory into more apps and actions.

Presenter cue

Do not disparage competitors. Acknowledge strengths, then name our three axes.

Transition

The final slide is the expansion thesis: what becomes possible once Echo works.

SLIDE 10 · 12:05–13:20

Echo is the beginning, not the boundary

One governed memory substrate can power an entire AI-native platform

> **Purpose** Close on inevitability and disciplined ambition.

Suggested speech

Echo begins with conversation truth. The shared agent turns that truth into reasoning. Signals turn reasoning into timely follow-through. Then future apps—projects, knowledge, service, operations—can reuse the same identity, memory, tools, connectors, and audit model.

That reuse is the platform advantage. A new app should contribute a new source of truth or a new governed action surface. It should not create a second identity system, a second memory store, or an unbounded agent.

We are building the AI people can trust with the work that matters: present before the prompt, natural in Persian, evidence-backed, and powerful without becoming omnipotent.

Capture what happened. Explain what it means. Act only with authority. Learn across time. That is NeurAI Platform—and Echo is how it begins.

Presenter cue

Finish slowly. Leave the four-sentence promise on screen for questions.

Transition

Invite questions on the wedge, the trust model, or the expansion sequence.

Q&A BACK POCKET

Short answers to likely questions

Question	Answer spine
Why not begin with a generic assistant?	Without a governed truth source, the assistant is generic and context has to be re-entered. Echo creates the first compounding memory asset.
Why not use one vendor for all AI?	Speech and reasoning have different quality/cost/failure profiles. The adapters keep providers replaceable while provenance makes the chosen lane visible.
Can the agent delete data?	No. The agent database role has no DELETE. Soft-delete UI actions are surface tools under the caller's authority; permanent purge is a named administrative path.
What makes Persian-first defensible?	It spans capture, UI direction, normalization, search, mixed-language transcription, names/glossary, digits/dates, typography, and acceptance fixtures—not a translation file.
What is shipped versus planned?	The repository contains the Echo and shared platform surfaces, routes, schema, queues, agent runtime, connectors and admin paths. Expansion apps are a disciplined thesis, not claimed as shipped.
Why Supabase/Postgres?	Identity, durable data, storage, queueing, RLS and grants can form one reviewable data plane. The database can enforce the wall even when application code is wrong.